SAMPLE

Building the Telecommunications of Tomorrow

The Elimination of Distance

Aiman JoeHanz

Free sample: Introduction & Chapter 1

putrabooks.github.io

Contents

Introduction: The Elimination of Distance

Chapter 1: The Signal and the Channel

Chapter 2: The Three Media

Chapter 3: The Switch and the Packet

Chapter 4: The Contract of Layers

Chapter 5: The Access Network

Chapter 6: The Cellular Decades

Chapter 7: The Long Haul

Chapter 8: Capacity and the Honest Limits

Chapter 9: Trust, Interception, and Survival

Chapter 10: The Money and the Decade Ahead

Connect

About this Series

This free sample contains the Introduction and Chapter 1. The complete book contains the full contents above.

Introduction: The Elimination of Distance

Telecommunications is the discipline of moving information without moving matter, and this volume earns its place on the shelf the way every volume in this series earns one: by refusing the marketing sentence and building the technology from its parts upward, until the reader can see the seams where it works and the seams where it fails. A phone call, a page load, a satellite pass overhead - each is a chain of physical events, negotiated contracts, and economic arrangements, and the whole industry is the accumulated discipline of keeping that chain intact. The chapters are built for two readings: one straight through, to acquire the argument, and a hundred partial ones afterward, opened at need before a purchase, a build, an audit, or an outage.

A telecommunications network is a machine whose purpose is to defeat distance. That sentence sounds like poetry and is in fact an engineering specification. Every element of the craft - the transmitter, the antenna, the fiber strand, the switch, the router, the protocol stack - exists to move a message from one point to another faster, cheaper, and more faithfully than a courier could, and each element is judged by exactly those three measures: speed, cost, and fidelity. Where the measures conflict, the conflicts are not incidental; they are the permanent weather of the industry, and every standard, every spectrum license, and every pricing plan is a negotiated settlement among them.

The settlement has physics as its referee. Information is not transported for free: it rides on waves that spread, attenuate, interfere, and fade, and there is a hard mathematical ceiling - a law of the channel, named for the engineer who first proved it - on how much information a given slice of bandwidth can carry through a given amount of noise. That ceiling is not a business problem and cannot be lobbied away; it is the reason the industry spends so much of its ingenuity on coding tricks, antenna geometries, and the patient economics of building ever-denser networks. The reader who holds the ceiling in mind can predict most of the industry's behavior without reading a single press release.

The claim this book makes is narrow and checkable. A network is a stack of contracts over a set of physical media, run by institutions that recover their costs under a license regime that decides who may transmit what, where. Understand the physics of the signal, the engineering of the media, the grammar of the switching, the legal skeleton of the spectrum, and the economics of the build - and no product announcement, outage report, or merger press release will be opaque again. That is the whole promise, and the series holds that structure outlives every product cycle.

The chapters walk a deliberate route. First the signal itself: what a wave is, what spectrum means, and how a message is stamped onto a carrier and pulled back off at the far end. Then the three media the industry actually builds with - the copper pair, the glass strand, and the air - and what each forgives and what each refuses. Then the switch: how an operator's cord panel became a packet forwarder, and why the packet won. Then the contract of layers: the protocol stack, the standards bodies, and the reason every layer leaks a little.

Then the two networks most readers touch daily - the access network, the last mile that economics makes hard, and the cellular network, whose generational ladder is the industry's loudest marketing and quietest engineering. After the anatomy comes the geography: the long haul of backbones, submarine cables, and the satellite constellations that promise the sky as infrastructure. Then the honest limits: the channel's mathematical ceiling, the difference between bandwidth and latency, and the queues that turn raw capacity into experienced service. Then trust and survival: interception, encryption, outage, and the restoration disciplines that keep a civilization's nervous system alive.

The book closes with the economics of the decade - who pays for infrastructure, how competition and subsidy shape the map, and what the constellation era changes and does not - and an operating program a reader can actually run.

A note on what deliberately does not appear. There are no vendor names, no carrier names, no generation trademarks, no market-share figures, and no forecasts dressed as facts. Where a number appears in the chapters, it is illustrative - a scale of magnitude, an order of cost - or structural, a property of the design itself, and the text says which. The cellular generations are named by their ordinals and their engineering, not by their logos.

Telecommunications is unusually rich in disputes, so the reader should begin with the three instruments the book sharpens and carries. The first is the honest reading of numbers: every rate, coverage claim, and availability figure in this industry is a measurement under conditions, and the conditions decide more than the figure - the book teaches the conditions and the subtractions, so the reader can price any claim rather than merely receive it. The second is the map habit: the network is a physical thing with a geography - sites, cables, landing stations, gateways, choke points - and every argument about speed, cost, or risk becomes tractable when placed on the map it is actually about. The third is the treaty lens: spectrum licenses, peering settlements, standards specifications, and interconnection contracts are the industry's real machinery of change, and the reader who reads them as treaties - with parties,

enforcement, and expiry - will predict the industry's behavior better than the reader who reads its launches.

One honest warning belongs at the door, because the subject invites its own mystification. Telecommunications is unusually rich in romantic framing - the shrinking world, the death of distance, the connected planet - and each phrase contains a real structural fact while omitting its costs. Distance has not died; it has been repriced, and the new price is this book's subject. The world has not shrunk; it has been wired, and the wiring has an anatomy, an economics, and a failure catalog that the marketing sentences omit. The reader who finishes this book should be able to take any romantic framing of the industry apart into its structural facts in a few sentences - and that ability, more than any fact, is what the series means by structural literacy.

The chapters are ordered as one continuous ascent and descent: the first half climbs the physical machine - the signal, the media, the switch, the stack - and the second half turns to the deployed world - the access network, the cellular decades, the long haul, the limits, the adversary, and the money. A reader new to the subject should take the chapters in order, because every chapter spends trust it earned in the one before. A reader with a specific decision pending may enter directly: the buyer auditing a connectivity vendor will find the boundary doctrine in the access, cellular, and economics chapters; the engineer sizing a link will find the load-bearing sections in the signal, media, and limits chapters; and the manager or official weighing an infrastructure program should read the economics chapter first, then the access and long-haul chapters again, more slowly. Each chapter closes with a single page that compresses its argument to keepable size, and the reader who retains only the ten closing pages retains the book.

Chapter 1: The Signal and the Channel

Every telecommunications system, from a ship's semaphore to a satellite downlink, reduces to the same three actors: a transmitter that stamps information onto a physical phenomenon, a channel that carries the phenomenon while degrading it, and a receiver that pulls the information back off and bets on having guessed right. This chapter builds those actors from first principles, because everything the industry sells - the tariff plan, the fiber build, the constellation launch - is a downstream consequence of what the channel permits and refuses.

An electromagnetic wave is a self-propagating handshake between an electric field and a magnetic field: a disturbance in one induces the other, and the pair travels through vacuum at a speed that is fixed by the universe rather than by any engineering decision. The wave has two properties the industry cares about above all. The first is frequency - how many times per second the field oscillates - and frequency is destiny: it determines the range of the wave, its tolerance for obstacles, how much information it can carry, and how much of it is legally available. The second is amplitude - the strength of the oscillation - and amplitude is budget: it costs transmit power, it fades with distance, and it is the first casualty of interference.

The full span of possible oscillation rates is called the spectrum, and it runs from waves that oscillate a few thousand times per second - used for the humble long-wave broadcasts that follow the curve of the Earth - to waves that oscillate trillions of times per second, which are light itself, and beyond into radiations that no civilian network uses. The spectrum is not a metaphor; it is a filing cabinet with an infinite number of drawers, of which only a modest fraction are useful for civilian communication, and each drawer is owned, leased, or contested.

Radio spectrum below roughly a few gigahertz bends around hills and penetrates walls, which is why it is the contested real estate of broadcasting and mobile service. Spectrum above that range behaves more like light: it travels in straighter lines, carries more per second, and is stopped by rain, leaves, and drywall. Neither band is better; they are different tools, and a large part of network engineering is matching the band to the job - coverage from the low drawers, capacity from the high ones.

Two waves of different frequency can share the same air without destroying each other, provided their oscillation rates are separated by more than the mess a receiver can untangle. This single fact - that frequency is a dimension along which users can be stacked - is the entire basis of spectrum

management. It is also why a receiver begins its work with a filter: a circuit tuned to pass one drawer of the filing cabinet and reject the rest, so that the receiver's job is to understand one voice in a room that is technically full of shouting.

A pure wave carries no information at all. A perfectly steady tone, transmitted forever, tells the receiver only that the transmitter is on; the receiver learns nothing new as the seconds pass. Information is change, and modulation is the craft of changing a wave in ways the transmitter controls and the receiver can measure. The wave offers three handles to grab: its amplitude can be nudged up and down, its frequency can be nudged up and down, or the timing of its oscillations - its phase - can be nudged ahead and behind. Each nudge encodes a symbol, and a symbol is the smallest unit of transmitted meaning: a pick from a fixed menu of possibilities agreed in advance.

The engineering trade sits in how big the menu is. The simplest scheme uses a menu of two symbols - a threshold nudged one way means one thing, the other way means the other - and the receiver's job is a coin flip with noise stacked against it. Bigger menus carry more bits per symbol, and the bits per symbol is one of the two dials that set a link's raw speed. The other dial is symbol rate - how many symbols per second the channel is asked to swallow. A channel's bandwidth, meaning the width of the slice of spectrum it is allowed to occupy, caps the symbol rate: the wider the slice, the faster the symbols can follow one another and still be told apart at the far end.

These two dials - bits per symbol and symbols per second - multiply into the raw data rate, and every generation of equipment in the industry's history is, at bottom, a new trick for turning one or both dials further than the last generation could.

The tricks have names worth knowing even in a vendor-neutral book, because the names describe structure rather than products. Quadrature schemes read two dials at once, treating amplitude and phase as a two-dimensional menu, which doubles the choices without doubling the required spectrum. Spread-spectrum schemes smear a narrow message across a wide slice of spectrum, deliberately trading bandwidth - which they waste cheerfully - for interference tolerance and the ability for many users to share one slice under a shared code discipline. Orthogonal frequency schemes chop a wide, brutal channel into many narrow, gentle subchannels and send them in parallel, which is how modern systems tame the frequency-selective cruelty of air. Each scheme is a settlement in the same negotiation: how much information per symbol, how much spectrum per second, how much noise tolerance the design buys back.

Between the transmitter's antenna and the receiver's filter, the world does its

work, and the work is hostile. The insult every channel delivers first is attenuation: the wave spreads its fixed transmit power over an ever-wider sphere, so the power arriving at a receiver falls with the square of the distance. Double the range and the received power drops to a quarter; this is the free-space law, and it is the reason transmit power and antenna gain are the first things an engineer budgets and the first things a regulator caps.

The insults compound. A wave that reflects off a building, a lake, or the ionosphere arrives as several copies of itself, slightly delayed from one another; the copies interfere, and the interference turns some frequencies within the slice into dead zones - a phenomenon called fading that can silence a link while the signal meter still reads strong. Motion adds insult to the insult: when either end moves, the multiple copies shift in relative phase, and the fade pattern slides across the frequencies like sunlight across a floor, so a mobile link must be designed to survive the constant possibility that its chosen frequency is momentarily dead.

Interference adds a different order of misery: another transmitter on or near the same slice, or the accumulated hiss of the thermal agitation of the world's electrons, raises a noise floor against which the received signal must stand out. The ratio of signal power to that floor is the channel's most-watched single number, and everything the receiver's modems do - the coding, the equalization, the antenna diversity - is a campaign to improve it or to survive its bad days.

The honest engineering response to all this is not to deny the insults but to add structure that absorbs them. Error-correcting codes append redundancy to every transmission - extra symbols that carry no new information but make the message reconstructible even after some symbols are lost. Modern codes approach the point where a message can be recovered from a signal that is, moment to moment, weaker than the noise itself, which sounds like a magic trick and is in fact careful mathematics: the redundancy is spread so finely across the transmission that the noise's random damage cannot erase enough of it at once.

Interleaving scrambles the message across time so that a burst of interference - a truck's ignition, a lightning crack - smears its damage across many recoverable fragments instead of punching one fatal hole. Retransmission protocols finish the job: when a block arrives damaged past repair, the receiver says so and the transmitter sends it again, turning an unreliable channel into a reliable service by spending time, which is the one currency the channel cannot manufacture.

All of these tricks buy their gains against a hard limit, and the limit is worth a

paragraph of its own, because the industry's economics hang from it. For any channel - a slice of spectrum, a fiber core, a copper pair - there is a maximum rate at which information can be carried with arbitrarily small error, and the rate is fixed by exactly two quantities: the width of the slice in cycles per second, and the ratio of signal power to noise power. Widen the slice or clean up the noise and the ceiling rises; nothing else - no pricing plan, no protocol, no antenna patent - moves it.

The law is the channel's constitution, and the practical consequence is that every network's capacity story is a story about acquiring more bandwidth, raising signal quality, or slicing the existing ceiling more cleverly among users. When a network operator promises speed, the promise is a wager on one of those three, and the reader who knows the constitution can price the wager.

The ceiling also explains the industry's structural trajectory. Because the ceiling scales with bandwidth, and bandwidth in the useful radio drawers is scarce while bandwidth in the optical drawers is enormous, the industry's history is a long migration of capacity from air to glass, with the air reserved for the one job glass cannot do: reaching things that move. The next chapter takes up the three media in turn - the copper pair, the glass strand, and the air

- and reads each as a material with a temperament: what it forgives, what it refuses, and what it costs to lay.

Between the modem's mathematics and the channel's insults stands the antenna, the least glamorous and most consequential component in the entire chain. An antenna is a structure that converts an electrical current into a radiating field, or the reverse, and its geometry decides where the energy goes. The simplest antennas radiate in all directions, which is honest generosity and mostly waste: a link between two fixed points does not need the sky's share of the power, and a base station does not need to warm the hillside behind it. Directional antennas concentrate the radiation toward the intended receiver, and the concentration factor - the gain - acts exactly like an amplifier without consuming a watt: the same transmit power, focused, arrives at the receiver multiplied.

Gain is bought with geometry and paid in narrowness, because a focused beam must be aimed, and the aiming tolerances of a long link are why dish feeds sit in rigid mounts and why a windstorm can take down a microwave link without breaking any of its parts.

Antenna engineering also decides the multipath story. A single antenna at the receiver sums every arriving copy of the signal - the direct path and every reflection - as one electrical mess, and the summation is arithmetic: copies in

phase reinforce, copies out of phase cancel, and the arithmetic lands wherever the geometry of the environment puts it. Two antennas, separated by even a fraction of a wavelength, sample the mess at different points, and the fade that kills one antenna's copy spares the other's - the diversity principle that puts two antennas on a wireless card and lets the receiver combine or select the stronger sample.

The modern extension is the antenna array: many small elements whose signals are combined with deliberate phase offsets, so that the combination itself steers the effective beam - the principle behind the beam-steering radios of the newest generations, which aim their energy at each scheduled user in turn rather than radiating a sector-wide glow. The reader should hold the array as the cellular chapter's key technology previewed here: steering the beam is the same as adding gain, and gain is capacity without new spectrum, which is why arrays became the industry's favorite trick of the era.

Every link is designed by a ledger, and the ledger deserves the reader's acquaintance because it is the honest form in which every radio engineering decision is actually made. The link budget starts with transmit power - a regulator-capped number - and walks the energy through its losses and gains:

the transmit antenna's gain, the free-space spreading loss computed from the distance, the media losses of foliage, walls, and rain, the receiver antenna's gain, and the receiver's own sensitivity floor, below which demodulation fails. Add the items and the margin is what remains: the difference between the power that arrives and the power the receiver needs, expressed in decibels, and every honest design carries margin for the fades the chapter listed, because a link with zero margin works on the day of installation and fails in the first thunderstorm.

The budget explains more engineering behavior than any marketing document. It explains why raising the data rate shortens the range: higher-order modulation demands a cleaner signal, so the sensitivity floor rises, the margin falls, and the working distance shrinks - which is why a wireless service's faster tier often covers less area than its slower tier. It explains why the millimeter drawers are short-range darlings: the budget's losses are larger, the fade margins must be fatter, and the cells shrink correspondingly. It explains why satellite terminals need dishes and handset radios do not: the satellite's enormous distance multiplies the spreading loss, so the budget must be balanced with gain on both ends.

And it explains, in one line, the industry's entire density campaign: when the spectrum is exhausted, the remaining lever is shrinking distances, because halving the distance buys back the margin that a wider drawer would have. A

reader who can sketch the budget can read a coverage map, a tariff tier, and a constellation's service promise with the same instrument.

The chapter closes its physics with the legal, because the drawer analogy was not decoration. In practice, spectrum is divided by international coordination into bands allocated to services - broadcast, mobile, fixed, satellite, science - with the coordination done at treaty conferences where nations align their allocations so equipment can roam and markets can scale. Below the international layer sits the national layer, where a regulator assigns the bands to users, by auction where the band is commercially contested, by license to fixed links where the geometry is exclusive, and - notably - by an unlicensed commons where the regulator deliberately holds a park open.

The unlicensed parks, where the cordless phone, the home wireless router, and a whole industry of short-range gadgets live, operate on a courtesy doctrine: low power, shared suffering, and etiquette built into the equipment, so that the park degrades gracefully under crowding instead of requiring permits. The reader should note the structural lesson of the parks: they are the industry's proof that spectrum scarcity is partly a property of transmit power and politeness rather than of physics alone, and they are the origin of several of the sharing etiquettes the cellular chapters treat as modern inventions.

A word on interference, the scarcity's shadow. Two systems sharing a drawer do not jam each other politely by turns; they raise each other's noise floors, and the floors rise invisibly - a link that worked for years degrades without any single assignable cause. The regulatory answer is a separation discipline: geographic spacing, power caps, out-of-band emission limits, and coordination procedures for the fixed links that must coexist. The engineering answer is the receiver's filter and the etiquette's backs-off, and the honest observation for the auditor is that interference disputes are decided by measurements and models rather than by rights, because the interference lives in the receiver's noise floor, which no property line encloses. The reader who holds that the spectrum is a shared physical environment - governed like one, disputed like one, and degraded like one - holds the correct frame for every spectrum controversy the decade will produce.

The receiver deserves its anatomy before the chapter closes, because it is where most of the discipline's mathematics actually runs. Its first stage is the filter, the tuned gate that passes the wanted drawer and rejects the neighbors. Its second is the amplifier, which raises the filtered signal above the electronics' own thermal noise - and the noise figure of that first stage is the receiver's engineering signature, because every decibel of noise added at the front door is a decibel subtracted from the link budget forever. Its third is the

equalizer, the adaptive filter that undoes the channel's frequency-selective damage by measuring the channel's response on training sequences the transmitter inserts at intervals, and applying the inverse response to what arrives - a piece of real-time mathematics that stands between the crude channel and the clean symbol.

Its fourth is the demodulator, the decision machinery that maps each measured symbol to its most probable meaning, and the modern demodulator is a probability engine: it soft-decides, keeping likelihoods rather than hard guesses, because the error-correcting decoder downstream works far better on probabilities than on coin flips. The modern receivers do all four stages in sampled digital arithmetic, and the honest observation is that the receiver, not the transmitter, is where a generation of equipment's performance is won: two systems with identical spectrum and power differ by their receivers, and the industry's quietest engineering competition has always run in the demodulator's silicon.

The training sequence is worth a sentence all to itself - it is the industry's most elegant habit: the transmitter deliberately spends some of every transmission's capacity sending a known pattern, which sounds like waste and is measurement. The receiver uses the known pattern to observe what the channel did to it, adapts its equalizer in real time, and spends the rest of the transmission benefiting from the knowledge. The reader should file the pattern as the discipline's general trade: capacity spent on measurement is capacity that buys back more than it costs, and the same trade reappears everywhere - the pilot carriers of the orthogonal schemes, the reference signals of the cellular generations, and the management traffic of the network's control planes.

The chapter has used decibels twice without introducing them, and the honest reader deserves the introduction, because the decibel is the telecommunications discipline's native arithmetic and every document the reader will ever audit uses it. A decibel expresses a ratio logarithmically: a factor of ten in power is ten decibels, a factor of two is three, a factor of one hundred is twenty. The logarithm converts the discipline's enormous ranges - the powers, losses, and gains that span a billionfold - into numbers a human can add and subtract in a ledger line, and the addition is the entire point: the link budget's items add because the logarithms add, which is why the budget is a one-line sum and why every radio engineer carries the small conversions - three for two, ten for ten - as mental furniture.

The reader who spends one hour internalizing the decibel reads spectrum regulations, coverage maps, and equipment datasheets fluently thereafter; the

reader who does not will find every subsequent audit in this industry conducted in a foreign number system. The honest caution is the same one every engineer learns: decibels express ratios, so every number in the discipline is relative to a reference, and the reader should always ask which - a power relative to a milliwatt, a bandwidth relative to a hertz, a gain relative to an ideal antenna - before comparing two numbers that look alike.

One machinery remains before the chapter closes, and it is the quietest: time. A receiver cannot demodulate what it cannot time-align with, so the transmitter embeds timing in its signal - a clock the receiver recovers and locks to - and the whole art of synchronization is the discipline of sharing time across a channel whose delay varies. The problem compounds up the stack: high-order symbol decisions require knowing when one symbol ends and the next begins, within a small fraction of a symbol; the phase-sensitive schemes require knowing where the carrier's cycle stands; and the networks of later chapters require their switches to agree on the time of day tightly enough to coordinate, to bill, and to authenticate.

The industry's answer is a hierarchy of shared time - receivers recover their timing from the signals themselves for the microseconds, networks sync their clocks over protocols for the milliseconds, and the whole structure ultimately anchors to the global timekeeping infrastructure the same way the naming hierarchy anchors to its roots. The honest observation for the reader is that synchronization is a dependency the network's users never see and the network's engineers budget like any other utility - and that its failures are exotic: the network that loses its shared time degrades in ways that look like interference, noise, or malice, and the diagnostic instinct that checks the clocks first is the mark of an engineer who has been burned.

Error correction earned one paragraph earlier and deserves its ledger, because the modern codes spend redundancy with a precision the older schemes could not approach. The classical codes guarded against the channel's average behavior: parity checks per block, with the receiver detecting what the parity could see and asking for retransmission of what it could not. The modern codes - the iterative decoders the industry calls by their inventors' initials - spread a message's information across the whole transmission so finely that the decoder can reconstruct from fragments, and they approach the channel's theoretical ceiling closely enough that the industry now speaks of codes in terms of their distance from the ceiling rather than their gains over the past.

The reader should hold the engineering trade in its modern form: redundancy is overhead, so a code's rate - the fraction of each transmission that is information - is a dial the network sets per link and per moment, deep

redundancy for the fading edge of the cell, shallow redundancy for the clean center, and the adaptive machinery of the cellular chapter turning the dial in real time. The honest summary is that the discipline's last thirty years bought their largest capacity gains not from new spectrum but from moving redundancy from a flat tax to a targeted one - and the same pattern, capacity recovered by removing waste rather than adding wealth, will recur throughout the book.

The interference story has subtler chapters the honest reader should know by name. Adjacent-channel interference arrives from a transmitter tuned next door in the filing cabinet, its energy slopping over the receiver's filter's imperfect edges - which is why the spectrum regulations police out-of-band emissions as strictly as in-band power, and why two neighboring licensees fight over filters as fiercely as over coverage. Intermodulation is the analog world's perversity: two clean transmitters mixing in some rusty joint or overloaded amplifier to produce spurious products at frequencies nobody licensed, the classic hunting ground of the interference investigator, whose instruments and patience are the industry's exorcists.

And the near-far problem is the receiver's own crisis: the strong nearby transmitter that saturates a receiver's front end while it strains to hear a weak distant one, the reason the cellular network commands its terminals' transmit powers continuously, turning every handset into a polite citizen of the air. The honest summary is that interference is never one phenomenon but a family of mechanisms, each with its own instruments and its own settlements, and the reader who hears the word in a dispute should ask which family member is on trial.

A closing note on the chapter's own scope: everything here belongs to the physical layer in the protocol chapter's vocabulary - the waves, the symbols, the channel's insults, and the receiver's mathematics - and the reader might reasonably ask where the physics ends and the network begins. The honest answer is that it ends at the frame: the first structure a receiver hands upward, a bounded run of symbols with a header and a checksum, is the atom of every networked conversation this book will build. The chapters that follow take the frame as their inheritance and never descend below it, because the division of labor the layering doctrine teaches is also this book's compact with the reader: each chapter builds on the last and hides the last's details.

And one honest sentence about the drawer below radio entirely: the wired channels of the next chapter obey the same mathematics - modulation, noise, ceiling - with a quiet medium instead of a noisy one, which is why the discipline calls the analysis common and the practice different. The reader who

holds the channel theory once holds it for glass and air alike, and the next chapter's media differences are differences of temperament, not of law.

The frequency reuse idea also previewed a discipline the cellular chapters will inherit: coordination. Two links on the same drawer in the same geography interfere unless someone coordinates them - a regulator with a license plan, an operator with a frequency plan, or an etiquette with a politeness rule - and the choice among the three coordinators is the deepest design decision a shared-medium network makes. The unlicensed parks of the spectrum section chose etiquette; the mobile estates chose licensing; the microwave routes chose case-by-case coordination - and the reader should hold that all three answers still work, each paying its coordination cost in a different currency: the politeness pays in capacity, the license pays in auction price, and the coordination pays in administrative delay.

A wave carries no information; modulation changes a carrier's amplitude, frequency, or phase, and each agreed change is a symbol. Bits per symbol times symbols per second is the raw rate; spectrum width caps the symbol rate. The channel attenuates with the square of distance, fades through multipath reflection, and interferes through neighbors and thermal noise; the receiver answers with filtering, coding, interleaving, and retransmission. Every channel has a mathematical ceiling set by bandwidth and signal-to-noise alone; capacity strategies are therefore always about more spectrum, cleaner signals, or smarter sharing, and a promised speed is a wager on one of the three. Radio's low drawers cover and penetrate; the high drawers carry and stop at walls; glass has the deepest drawer of all. Everything the industry builds is downstream of this physics.

End of free sample. The full book is available at putrabooks.github.io